

# Apache Spark Workshop

J.T. Halbert

May 6, 2015

# Workshop goals

- ▶ Work a data exploration and analysis problem with Spark.
- ▶ Learn enough Scala to appreciate the collections library.
- ▶ Leave feeling you know enough to begin teaching yourself.

# Workshop non-goals

- ▶ Deploy and monitor Spark Applications
- ▶ Graphs
- ▶ MLlib (well... maybe if we have time)

# Overview of resources

- ▶ This presentation and “answers” to the “exercises” are available at  
<https://github.com/jt-halbert/spark-workshop/>.
- ▶ The ETL code that is making our exploration of the Enron dataset so convenient is available at  
<https://github.com/medale/spark-mail/>.
- ▶ LARGE PORTIONS of this presentation are pulled from Markus' work. Thanks Markus!!
- ▶ The data we will need is already in the tar you downloaded (or the VM). It is a subset of the famous ENRON dataset.
  - ▶ <https://www.cs.cmu.edu/~./enron/>

# Who are we?

- ▶ I am Tetra's Chief Data Scientist and I help certain people learn certain things about certain parts of their data.
- ▶ Tetra Concepts, LLC is the finest collection of development talent anyone could ask for (AND SOME OF THEM ARE WALKING AMONG YOU!!!!11!!1!!!!)

# Why Apache Spark?

- ▶ Excellent question

# Data Science is a filthy job

- ▶ I am not even sure what Data Science is. They put Science right in the name, so it must be pretty serious right?
- ▶ I like to think it is the disciplined application of a scientific mindset to that nebulous thing called “data.”
- ▶ In my experience it is three activities

# The Three Big Things

1. Find a Problem.
2. Find a Solution.
3. Automate?

Be very careful about the order.



# Why Apache Spark?

- ▶ Spark gives you a way to explore small, medium, large, (very large ?) data in a convenient way.
  - ▶ You can actually explore distributed datasets: lazy evaluation and a rich collections api.
  - ▶ You can scale your exploratory code up to a full job relatively quickly: REPL driven development.
- ▶ It wraps an increasing amount of the Hadoop Ecosystem and plays naturally.

Your customer wants pretty little magical things



Figure 1: Spark wraps a lot of other people's toys

# Let's get started

- ▶ The first step is to learn enough Scala to be dangerous.

# Combinator functions on Scala collections

- ▶ Examples: map, flatMap, filter, reduce, fold, aggregate
- ▶ Background - Combinatory logic, higher-order functions...

# Combinatory Logic

Moses Schönfinkel and Haskell Curry in the 1920s

*[C]ombinator is a higher-order function that uses only function application and earlier defined combinators to define a result from its arguments [Combinatory Logic @wikipedia\_combinatory\_2014]*

A *Higher-Order Function* is a function that takes functions as arguments or returns function.

# Functional Programming

- ▶ An approach/style of programming that deals with expressions and values rather than statements.
- ▶ Functions are treated (to the extent possible) as Mathematical Functions (no side effects, deterministic)

# map

- ▶ Applies a given function to every element of a collection
- ▶ Returns collection of outputs of that function
- ▶ input argument - same type as collection type
- ▶ return type - can be any type

## Map




Figure 2:



## map - Scala

```
def computeLength(w: String): Int = w.length

val words = List("when", "shall", "we", "three",
  "meet", "again")
val lengths = words.map(computeLength)

> lengths    : List[Int] = List(4, 5, 2, 5, 4, 5)
```

## map - Scala syntactic sugar

```
//anonymous function (specifying input arg type)  
val list2 = words.map((w: String) => w.length)
```

```
//let compiler infer arguments type  
val list3 = words.map(w => w.length)
```

```
//use positionally matched argument  
val list4 = words.map(_.length)
```

# map - ScalaDoc

See immutable List ScalaDoc

List[+A]

...

```
final def map[B](f: (A) => B): List[B]
```

- ▶ Builds a new collection by applying a function to all elements of this list.
- ▶ B - the element type of the returned collection.
- ▶ f - the function to apply to each element.
- ▶ returns - a new list resulting from applying the given function f to each element of this list and collecting the results.

# flatMap

- ▶ ScalaDoc:

```
List[+A]
```

```
...
```

```
def flatMap[B](f: (A) =>  
    GenTraversableOnce[B]): List[B]
```

- ▶ GenTraversableOnce - List, Array, Option...
- ▶ can be empty collection or None
- ▶ flatMap takes each element in the GenTraversableOnce and puts it in order to output List[B]
- ▶ removes inner nesting - flattens
- ▶ output list can be smaller or empty (if intermediates were empty)

## What If?

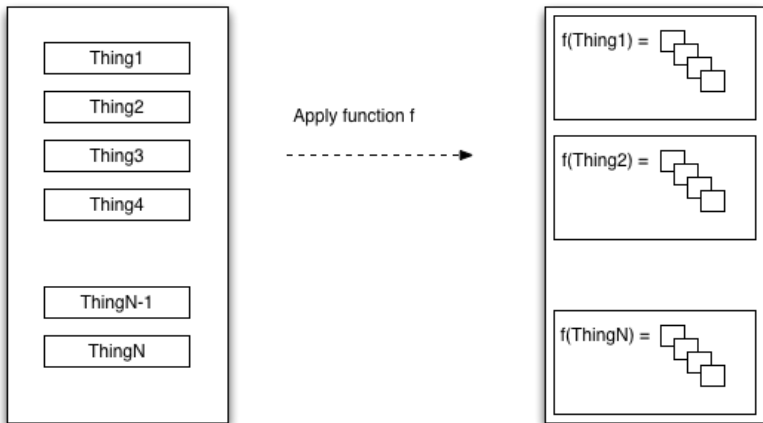


Figure 3:

## flatMap

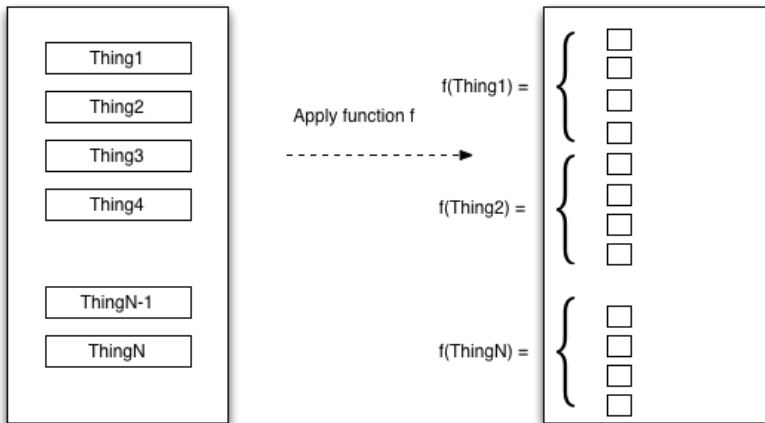


Figure 4:

## flatMap Example

```
val macbeth = """When shall we three meet again?  
|In thunder, lightning, or in rain?""".stripMargin  
val macLines = macbeth.split("\n")  
// macLines: Array[String] = Array(  
//   When shall we three meet again?,  
//   In thunder, lightning, or in rain?)  
  
//Non-word character split  
val macWordsNested: Array[Array[String]] =  
    macLines.map{line => line.split("""\W+""")}  
//Array(Array(When, shall, we, three, meet, again),  
//      Array(In, thunder, lightning, or, in, rain))  
  
val macWords: Array[String] =  
    macLines.flatMap{line => line.split("""\W+""")}  
//Array(When, shall, we, three, meet, again, In,  
//      thunder, lightning, or, in, rain)
```

# filter

```
List[+A]
```

```
...
```

```
def filter(p: (A) => Boolean): List[A]
```

- ▶ selects all elements of this list which satisfy a predicate.
- ▶ returns - a new list consisting of all elements of this list that satisfy the given predicate p. The order of the elements is preserved.



## filter Example

```
val macWordsLower = macWords.map{_.toLowerCase}  
//Array(when, shall, we, three, meet, again, in, thunder,  
//      lightning, or, in, rain)  
  
val stopWords = List("in","it","let","no","or","the")  
val withoutStopWords =  
    macWordsLower.filter(word => !stopWords.contains(word))  
// Array(when, shall, we, three, meet, again, thunder,  
//      lightning, rain)
```

# reduce

```
List[+A]
```

```
...
```

```
def reduce[A1 >: A](op: (A1, A1) => A1): A1
```

- ▶ Creates one cumulative value using the specified associative binary operator.
- ▶ A1 - A type parameter for the binary operator, a supertype (super or same) of A. (List is covariant +A)
- ▶ op - A binary operator that must be associative.
- ▶ returns - The result of applying op between all the elements if the list is nonempty. Result is same type as (or supertype of) list type.
- ▶ UnsupportedOperationException if this list is empty.

## Reduce

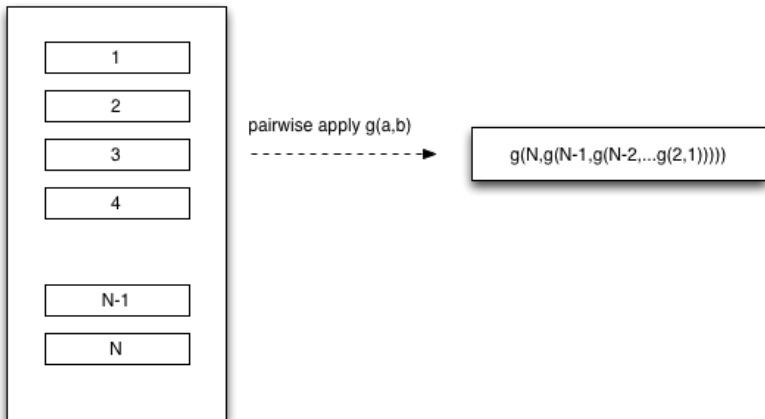


Figure 5:

## reduce Example

```
//beware of overflow if using default Int!  
val numberOfAttachments: List[Long] =  
    List(0, 3, 4, 1, 5)  
val totalAttachments =  
    numberOfAttachments.reduce((x, y) => x + y)  
//Order unspecified/non-deterministic, but one  
//execution could be:  
//0 + 3 = 3, 3 + 4 = 7,  
//7 + 1 = 8, 8 + 5 = 13  
  
val emptyList: List[Long] = Nil  
//UnsupportedOperationException  
emptyList.reduce((x, y) => x + y)
```

# fold

List[+A]

...

```
def fold[A1 >: A](z: A1)(op: (A1, A1) => A1): A1
```

- ▶ Very similar to reduce but takes start value z (a neutral value, e.g. 0 for addition, 1 for multiplication, Nil for list concatenation)
- ▶ returns start value z for empty list
- ▶ Note: See also foldLeft/Right (return completely different type)

```
foldLeft[B](z: B)(f: (B, A) => B): B
```

## fold Example

```
val numbers = List(1, 4, 5, 7, 8, 11)
val evenCount = numbers.fold(0) { (count, currVal) =>
    println(s"Count: $count, value: $currVal")
    if (currVal % 2 == 0) {
        count + 1
    } else {
        count
    }
}
```

Count: 0, value: 1

Count: 0, value: 4

Count: 1, value: 5

Count: 1, value: 7

Count: 1, value: 8

Count: 2, value: 11

evenCount: Int = 2

## So what does this have to do with Apache Spark?

- ▶ Resilient Distributed Dataset (RDD)
- ▶ From API docs: “immutable, partitioned collection of elements that can be operated on in parallel”
- ▶ map, flatMap, filter, reduce, fold, aggregate...

# Spark - RDD API

- ▶ RDD API
- ▶ Transforms - map, flatMap, filter, reduce, fold, aggregate...
  - ▶ Lazy evaluation (not evaluated until action!)
- ▶ Actions - count, collect, first, take, saveAsTextFile...



## Spark - From RDD to PairRDDFunctions

- ▶ If an RDD contains tuples (K,V) - can apply PairRDDFunctions
- ▶ Uses implicit conversion of RDD to PairRDDFunctions
- ▶ In 1.2 and previous versions available by importing `org.apache.spark.SparkContext._`

From `org.apache.spark.SparkContext`:

```
implicit def rddToPairRDDFunctions[K, V](  
  rdd: RDD[(K, V)])  
  (implicit kt: ClassTag[K],  
   vt: ClassTag[V],  
   ord: Ordering[K] = null) = {  
    new PairRDDFunctions(rdd)  
  }
```

# PairRDDFunctions

- ▶ keys, values - return RDD of keys/values
- ▶ mapValues - transform each value with a given function
- ▶ flatMapValues - flatMap each value (0, 1 or more output per value)
- ▶ groupByKey - `RDD[(K, Iterable[V])]`
  - ▶ Note: expensive for aggregation/sum - use `reduce/aggregateByKey`!
- ▶ reduceByKey - return same type as value type
- ▶ foldByKey - zero/neutral starting value
- ▶ aggregateByKey - can return different type
- ▶ join (left/rightOuterJoin), cogroup ...

# From RDD to DoubleRDDFunctions

- ▶ From API docs: “Extra functions available on RDDs of Doubles through an implicit conversion. Import `org.apache.spark.SparkContext._`”

From `org.apache.spark.SparkContext`:

```
implicit def doubleRDDToDoubleRDDFunctions(  
  rdd: RDD[Double])  
  = new DoubleRDDFunctions(rdd)
```

# DoubleRDDFunctions

- ▶ mean, stddev, stats (count, mean, stddev, min, max)
- ▶ sum
- ▶ histogram

# Example 1 - Mail Folder Statistics In MapReduce

- ▶ What are the least/most/average number of folders per user?
- ▶ Each MailRecord has user name and folder name

```
lay-k/      <- mailFields(UserName)
  business  <- mailFields(FolderName)
  family
  enron
  inbox
  ...
```

# Hadoop Mail Folder Stats - Mapper

- ▶ read each mail record
- ▶ emits key: userName, value: folderName for each email

# Hadoop Mail Folder Stats - Reducer

- ▶ reduce method
  - ▶ create set from values for a given key (unique folder names per user)
  - ▶ `set.size ==` folder count
  - ▶ keep adding up all `set.size` (`totalNumberOfFolders`)
  - ▶ one up counter for each key (`totalUsers`)
  - ▶ keep track of min/max count
- ▶ cleanup method
  - ▶ compute average for this partition:  
`totalNumberOfFolders/totalUsers`
  - ▶ write out min, max, `totalNumberOfFolders`, `totalUsers`, `avgPerPartition`

# Hadoop Mail Folder Stats - Driver

- ▶ Set Input/OutputFormat
- ▶ Number of reducers



# Hadoop Mail Folder Stats - Results

- ▶ if only one reducer - results are overall lowest/highest/avg
- ▶ if multiple reducers
  - ▶ post-processing overall lowest/highest
  - ▶ add totalNumberOfFolders and totalUsers to compute overall average

# Hadoop Mapper

```
public void map(AvroKey<MailRecord> key,
NullWritable value, Context context) throws ... {
    MailRecord mailRecord = key.datum();
    Map<CharSequence, CharSequence> mailFields =
        mailRecord.getMailFields();
    CharSequence userName =
        mailFields.get(AvroMailMessageProcessor.USER_NAME);
    CharSequence folderName =
        mailFields.get(AvroMailMessageProcessor.FOLDER_NAME);
    userKey.set(userName.toString());
    folderValue.set(folderName.toString());
    context.write(userKey, folderValue);
}
```

# Hadoop Reducer

```
public void reduce(Text userKey,
    Iterable<Text> folderValues,
    Context context) throws ... {
    Set<String> uniqueFolders = new HashSet<String>();
    for (Text folder : folderValues) {
        uniqueFolders.add(folder.toString());
    }
    int count = uniqueFolder.size();
    if (count > maxCount) maxCount = count;
    if (count < minCount) minCount = count;
    totalNumberOfFolder += count
    totalUsers++
}
...
public void cleanup...
//write min, max, totalNumberOfFolders,
//totalUsers, avgPerPartition
```

Let's get to work

# Warm-up

- ▶ Who sent the most email?
- ▶ Hint: if you have an `RDD[(String, Int)]` you can do this:  
`myRdd.top(10)(Ordering.by(_._2)).foreach(println)`

# Problem Statement

- ▶ Build a Markov model of Vince's sentences.
  - ▶ We'll talk through what that is as we progress.

## Step 1

- ▶ Find all the email bodies from Vince that contain mostly letters.
  - ▶ This is somewhat arbitrary, but it limits to emails without spreadsheets pasted into them.
- ▶ Hint: Use this (if you want)

```
def ratioLetters(s: String): Double = {  
  val (tot, n) = s.replaceAll("""\s+""", "")  
                  .foldLeft((0.0, 0.0)) {  
    (acc, c) => {  
      (acc._1 + 1, acc._2 + (if (c.isLetter) 1 else 0))  
    }  
  }  
  n/tot  
}
```

## Step 2

- ▶ Get rid of text below sigline and break into sentences.
- ▶ Hint:

```
val uptoSig = ".*?\s+Vince".r
val sentenceSplitter = """(?<=[.!\?])\s+(?=[A-Z])"""
```

- ▶ Note that regexes have methods for working on strings, and vice versa



## Step 3

- ▶ Build an RDD of trigrams

- ▶ e.g. "This is a simple sentence." becomes

```
("","","This")  
("","This","is")  
("This","is","a")  
("is","a","simple")  
("a","simple","sentence.")  
("simple","sentence.","")  
("sentence.", "", "")
```

- ▶ Hint:

```
_.split("\\s+")
```

and

```
.sliding(3)
```

## Step 4 (Challenge)

- ▶ Build the model:
  - ▶ Reduce to `RDD[(List[String],(String, Int))]` of Antecedent Pair of Word followed by Consequent Word and count.

## Step 5 (Challenge)

- ▶ Build a Vince Sentence Generator.
- ▶ Hint: Multinomial Sampler

```
def sampleMultinomial[T](dist: List[(T, Int)]): T = {  
  val p = scala.util.Random.nextDouble*dist.map(_._2).sum  
  def recur(acc: Int, d: List[(T, Int)]): T = d match {  
    case (t,n) :: Nil => t  
    case (t,n) :: ds  => if (acc + n >= p) t else {  
      recur(acc+n,ds)  
    }  
  }  
  recur(0,dist)  
}
```